

# AI and Games Semester 2 Project Report

## Group 16

Evripidis Christodoulides, Jonathan Haylett, Wenchang Liu, Huan Yang

May 11, 2020

## Contents

<b>1</b>	<b>Task Overview</b>	<b>2</b>
1.1	Estimating the Reaction Function of the Follower . . . . .	2
1.2	Calculating the Price to Publish . . . . .	2
<b>2</b>	<b>Linear Regression with Online Learning</b>	<b>2</b>
2.1	Linear Regression . . . . .	2
2.2	Online Learning: The Moving Window Approach . . . . .	2
2.3	Algorithm Implementation . . . . .	3
2.4	Evaluation Based on Profits (Moving Windows) . . . . .	4
<b>3</b>	<b>Evaluation of Different Regression Models</b>	<b>4</b>
3.1	Evaluation of Linear Regression . . . . .	4
3.2	Evaluation of Decision Tree Regression . . . . .	5
3.3	Other Regression Experiments . . . . .	5
<b>4</b>	<b>Other Ideas – Explore Time Series</b>	<b>5</b>
4.1	Calculate slope between current day and previous day . . . . .	6
4.2	Sequential Model . . . . .	6

# 1 Task Overview

The objective of this project is to develop a leader AI program that tries to get as much profit as it can by publishing its price every day for 30 days of simulation in a Stackelberg game against 3 follower AI programs respectively. The process of achieving this objective is divided into 2 main steps: estimating the reaction function of the follower, and calculating the price to publish.

As mentioned in section 1.2, since the procedure after estimating the reaction function of the follower is the same, we will mainly discuss how to do the regression to estimate the reaction function more accurately. We have implemented a linear regression leader agent with an online learning algorithm that can make a greater profit than SimpleLeader. Also, we experimented with other regression models including decision trees, SVM, etc. Eventually, we propose some other ideas which may perform better for this task, but we did not have time to implement them.

## 1.1 Estimating the Reaction Function of the Follower

In order to determine the price we publish for each day, the leader will try to estimate the reaction of the follower to the price it publishes, i.e.  $U_F = F(U_L)$ . Theoretically, the more accurate our estimation is, the more profit we will get. The data that we have available to do the estimation with is the 100 pairs of historical data. This is a regression task where we need a mapping function that maps a price that the leader publishes to a price that the follower publishes. The idea is that if we can train the model to fit most of the historical data, it can make predictions that are close to what the follower will actually publish. Such models including linear regression and its variations.

## 1.2 Calculating the Price to Publish

After estimating the reaction function of the follower, we can then bring the reaction function into the demand model of the leader, and then we will receive a formula with only one parameter —  $U_L$ . In order to calculate the maximum value of this formula, we will have to differentiate it to find which  $U_L$  can make it reach the maximum value, and this  $U_L$  is what we will publish. Finally, we can calculate the daily profit. No matter what models we use for learning the reaction function, the process for calculating the price to publish is always the same.

# 2 Linear Regression with Online Learning

## 2.1 Linear Regression

The simplest regression idea is to use linear regression. The basic idea of it is that we assume the follower's reaction function is a linear function, i.e. say  $U_F = a \cdot U_L + b$ , where  $a$  and  $b$  are unknown parameters, and an approximation function  $U_F = a^* \cdot U_L + b^*$ , where  $a^*$  and  $b^*$  are the parameters we need to calculate. Then we use the square error sum at the given data to measure the closeness, and we update the parameters  $a^*$  and  $b^*$  using the historical data we have so that our predicted linear function  $U_F = a^* \cdot U_L + b^*$  can be as close to the original linear function as possible.

## 2.2 Online Learning: The Moving Window Approach

In addition to the 100 days of historical data, we still have new data generated during the simulation. The moving window approach is designed to make use of these pieces of information. For example, if we use the window size 100, then for each new day of the simulation, we will calculate the parameters using the previous 100 days' data starting from the current day. In other words, we will discard data that has reached a certain age and take an equal amount of new data into consideration. The advantage is that if the environment is continuously changing, we can discard outdated data and update our model with the latest data.

## 2.3 Algorithm Implementation

$$\hat{a}^* = \frac{\sum_{t=1}^T u_L^2(t) \sum_{t=1}^T u_F^R(t) - \sum_{t=1}^T u_L(t) \sum_{t=1}^T u_L(t) u_F^R(t)}{T \sum_{t=1}^T u_L^2(t) - \left( \sum_{t=1}^T u_L(t) \right)^2}$$

$$\hat{b}^* = \frac{T \sum_{t=1}^T u_L(t) u_F^R(t) - \sum_{t=1}^T u_L(t) \sum_{t=1}^T u_F^R(t)}{T \sum_{t=1}^T u_L^2(t) - \left( \sum_{t=1}^T u_L(t) \right)^2}$$

Figure 1: Formulae for calculating  $a^*$  and  $b^*$

For each new day, we will first calculate the parameters  $a^*$  and  $b^*$  according to the formulas in figure 1 using the same amount of data as the window size, then we put  $U_F = a^* \cdot U_L + b^*$  into the daily profit formula  $(U_L - 1) \cdot (2 - U_L + 0.3 \cdot U_F)$  and get  $(U_L - 1) \cdot (2 - U_L + 0.3 \cdot (a^* \cdot U_L + b^*))$ , i.e.  $(U_L - 1) \cdot (2 - U_L + 0.3 \cdot a^* \cdot U_L + 0.3 \cdot b^*)$ . We then differentiate this formula to get  $U_L \cdot (-2 + 0.6 \cdot b^*) + 3 + 0.3 \cdot a^* - 0.3 \cdot 0b^*$ , and let the differentiated formula equals 0, then we can solve the equation to get  $U_L = (-3 - 0.3 \cdot a^* + 0.3 \cdot b^*) / (-2 + 0.6 \cdot b^*)$  which will theoretically make the leader receive maximum profits. Eventually, we just put the calculated  $a^*$  and  $b^*$  into this formula to get the price we should publish. The flow diagram of the linear regression algorithm with the moving window approach is shown below in figure 2.

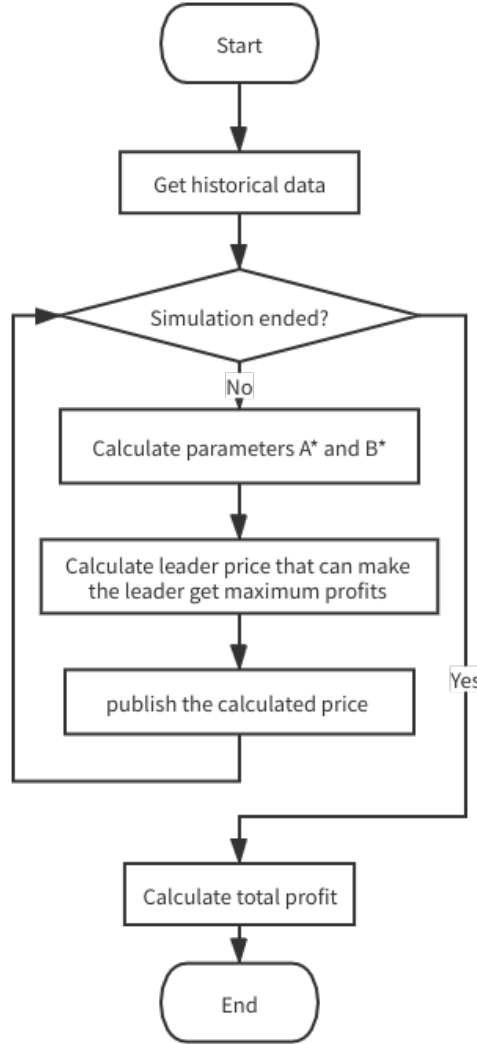


Figure 2: Flow Diagram of Linear Regression with Moving Window Algorithm

## 2.4 Evaluation Based on Profits (Moving Windows)

We tested our model with different window sizes against all followers, and we also compared our profits to the given SimpleLeader’s profits. The results are shown below in table 1.

Leader	Mk1	Mk2	Mk3
SimpleLeader	17.491040	16.897467	19.423810
Window Size 10	17.405104	16.932770	19.411194
Window Size 25	17.484606	16.953413	19.464302
Window Size 40	17.535341	16.95259	19.488014
Window Size 55	17.547180	16.952982	19.487911
Window Size 70	17.552155	16.953724	19.488184
Window Size 85	17.554714	16.953964	19.488335
Window Size 100	17.555769	16.955841	19.488330
Window Size 115	17.556887	16.956198	19.488293
Window Size 130	17.557184	16.956450	19.488283

Table 1: Comparison of Profits with Different Window Sizes

As we can see from the table, we made more profits by estimating the follower reaction function and publishing our price according to the reaction function than randomly publishing prices like the SimpleLeader. The best performance against Mk1 and Mk2 is achieved by using window size 130, and the best performance against Mk3 is achieved using window size 85.

## 3 Evaluation of Different Regression Models

Apart from implementing the linear regression in Java, we explored more regression models and evaluated them with R-squared scores with Scikit-learn. The R-squared score is a good way to evaluate the performance of a regression. The score of a regression can be ranged from 0 to 1, which is a relative evaluation. Such a score has an advantage over other absolute scores in that we cannot easily determine how high an absolute score can be regarded as a good score. However, it is safe to say an R-squared score close to 1 is a good one.

### 3.1 Evaluation of Linear Regression

We evaluated linear regression on mk1, mk2, mk3 historical data, and their R-squared scores are 0.19, 0.11, 0.72 respectively. As we know, an R-squared score below 0.4 means the regression is not that good, so our linear regressions are not good enough except for mk3.




Figure 3: Linear Regression on Mk1 Historical Data

In linear regression, we can imagine that we need to draw a straight line to fit most of the data points and make sure other data points are not too far away. But as we can see in figure 3, the data points from mk1 are scattered, so it is very hard to fit many points with a single line. This is why we want to experiment with other non-linear regression models and see if they can achieve better R-squared scores.

### 3.2 Evaluation of Decision Tree Regression

The decision tree regression can learn complex rule-based patterns; however, it seems to suffer from overfitting issues here according to our experiments. We split the historical data into a training set and testing set. While the decision tree regression model achieves a 0.93 R-squared score on the training set, it achieves a score of -1.58 on the testing set. A negative value means the model can be arbitrarily worse. As shown in figure 4, the model seems to find the wrong way to fit the data points.




Figure 4: Decision Tree Regression on Mk1 Historical Data

### 3.3 Other Regression Experiments

We have also tried other popular regression models including SVM, KNN, random forest, and AdaBoost on Mk1 and none of them worked better than linear regression. Table 2 shows the R-squared scores obtained from these experiments.

Model	Training R Score	Testing R Score
Linear	0.20563019704498864	0.12381456611660624
Decision Tree	0.9334617516195677	-1.5818682765183536
SVM	-0.003296979174702175	-0.17390958609995066
KNN	0.3395643813770154	0.06677303550060165
Random Forest	0.7072866944733289	-0.047113503710458415
AdaBoost	0.40155919508932547	0.05755037520545825

Table 2: Results of Experimentation with Different Regression Models on Mk1 Historical Data

## 4 Other Ideas – Explore Time Series

From the results of our experiments, we think it is unlikely that we will be able to make precise predictions by simply fitting all the historical points together. The reaction function might be affected by the time series. One possibility is that there are several sets of parameters for a follower based on the current date. Another possibility is that the follower is not only reacting to the leader simply according to the current published price of the leader, but also taking several

previous days' data into consideration. We have thought about several approaches to explore the time series problem.

#### 4.1 Calculate slope between current day and previous day

Assuming the follower has several different linear reaction functions, then each function will have a different slope. Therefore, theoretically, after we calculate the slope between the current day's price and the previous day's price  $\frac{(follower\_i - follower\_i-1)}{(leader\_i - leader\_i-1)}$  as y-axis and using dates as x-axis to plot the line graph, we may see the turning point of the slopes and that is when the parameters have changed. If we can chunk all the historical data into several parts, then we can do the same as section 2 to do a more accurate linear regression for each part separately.

#### 4.2 Sequential Model

If the follower decides its price by a sequence of data, e.g. previous 3 days' prices, then a sequential model might help. If we put all historical data together, it is like the bag of words model in NLP, which does not consider the sequences of data. Similarly, we can simulate the skip-gram idea by taking the leader's price from consecutive several days as the input to the regression model, and do a multivariable regression. Alternatively, we can directly try other sequential models such as RNN networks to predict a sequence of data.